

Prompt Types and Techniques


Module 2, Chapter 2

Kirsten Hunter

Now that you know what goes into a good prompt, let's talk about technique. There's more than one way to prompt an LLM, and different approaches unlock different capabilities. In this chapter we'll cover the most useful ones — and when to reach for each.

Zero-shot prompting

Just ask...



Zero-shot prompting is the default — you give the model a task with no examples, no setup, just the request. "Translate this to Spanish." "Summarize this paragraph." It works well for clear, simple tasks where the model already has strong training. Most people start here without knowing it has a name.

Few-shot prompting

Show before you tell


show before you tell



Few-shot prompting means giving the model two or three examples of what you want before your actual request. You're showing it the pattern. If you need responses in a very specific format — a particular tone, a consistent structure — few-shot is the fastest way to get there. The model pattern-matches against your examples rather than guessing.

Chain-of-thought prompting

Think it through



Adding "think step by step" to a prompt sounds almost too simple — but it meaningfully improves accuracy on anything involving logic, math, or multi-step reasoning. You're asking the model to show its work rather than jump to an answer. That process of working through steps catches errors that a direct answer would skip right past.

Role and persona prompting


Set the perspective



Telling the model who it is changes how it responds. "You are a senior software engineer reviewing a pull request" gets you a different answer than "you are a patient teacher explaining this to a beginner." The vocabulary shifts, the assumptions shift, the tone shifts. Role prompting is one of the fastest ways to calibrate output for a specific audience or purpose.

Instruction-based prompting

Direct, not conversational



Instruction-based prompting means writing explicit directives rather than conversational asks. Instead of "can you help me with..." you write "do this, in this format, for this audience." It feels less natural at first but produces more consistent results — especially in API contexts or system prompts where you need predictable behavior every time.

Iterative prompting

Your first prompt is a draft



Treat your first response as a starting point, not a final answer. Follow-up prompts like "make it shorter," "add a concrete example," or "rewrite this for a technical audience" let you steer incrementally rather than rewriting the whole prompt from scratch. Prompting is a conversation. The model holds the context of what it just produced and can refine from there.

Combining techniques

They stack!

"You are a customer support specialist. [Role]
Here are two example responses that match our
tone: [Few-shot] Now respond to this customer
complaint, step by step. [Chain-of-thought]"

These techniques aren't mutually exclusive. A role prompt combined with few-shot examples and a chain-of-thought instruction produces dramatically more consistent, targeted output than any single technique alone. Once you're comfortable with each one individually, start stacking them. The combinations are where the real control comes from.

Which technique when?


Think brief, not question

Situation	Technique
Simple, clear task	Zero-shot
Need consistent format =	Few-shot
Logic or reasoning task	Chain-of-thought
Audience or tone matters	Role prompting
Need predictable structure	Instruction-based
Response needs refinement	Iterative

Here's a quick framework. Simple task with no special format requirements — zero-shot is fine. Need consistency or a specific pattern — add examples. Anything involving reasoning or math — chain-of-thought. Audience matters — set a role. Need reliable structure every time — instruction-based. And whenever the first response is close but not quite there — iterate. You'll develop instincts for this quickly.

Key Takeaways

The toolkit



Zero-shot — just ask

Few-shot — show the pattern

Chain-of-thought — think step by step

Role prompting — set the perspective

Instruction-based — direct and explicit

Iterative — refine from the first response

Four elements. Task, context, format, constraints. You don't always need all four — but when a prompt isn't landing, one of them is almost always missing. Start there. In the next chapter we're going to build on this foundation with specific techniques — few-shot prompting, chain-of-thought, role assignment — and look at when to use each one.