
Module 3, Chapter 1

Quick Tips for Better Prompts

You've learned the framework. You know the five mistakes. Now here's the part that actually sticks: a handful of habits that separate people who get consistently useful outputs from people who keep rewriting the same prompt wondering why it isn't working.

None of these are complicated. Most take three extra seconds. All of them make a real difference.

Lead with the output, not the topic

The single highest-leverage change you can make is to start your prompt with what you want to receive, not what you want to talk about. Compare these two:

"Climate change is a complex topic. Can you help me write something about it for our newsletter?"

"Write a 250-word newsletter introduction about how our company is reducing its carbon footprint. Conversational tone, no jargon, ends with a call to action."

The second prompt leads with the deliverable. The model knows immediately what done looks like. Try it in Groq (or your preferred LLM) with a topic you're actually working on — the difference in first-draft quality is usually significant.

Use "as if" to calibrate complexity

One of the fastest ways to hit the right level of explanation is the "as if" construction: *"Explain this as if I'm a smart person who has never worked in finance"* or *"Write this as if it's for a senior engineer who doesn't need background."* It sets both the depth and the assumed knowledge in a single phrase, and it works reliably across models.

Give it a bad example to avoid

Most people give models examples of what they want. Fewer people tell the model what they don't want — and that's often more precise. *"Don't use corporate buzzwords like 'leverage' or*

'synergy'" or "Avoid bullet points — write in prose" are constraints that are much easier for a model to follow than vague instructions like "keep it natural."

Ask for options, not answers

When you're not sure what you want yet, don't ask for the thing — ask for three versions of the thing. *"Give me three different opening lines for this email, each with a different tone: professional, friendly, and direct."* You'll almost always find a direction you like in the options, or be able to combine elements from two of them. It's faster than iterating on a single response and gives you creative range you wouldn't have thought of alone.

Separate the thinking from the writing

For anything complex, ask the model to think first and write second. *"Before writing the email, briefly outline what points you plan to make, then write it."* This catches wrong assumptions before they're baked into a full draft — and produces noticeably better output on anything longer than a paragraph.

One more thing: save your best prompts

The prompts that work are worth keeping. A simple document — even a notes app — with your best prompts for recurring tasks (meeting summaries, email drafts, code review requests) means you're not starting from scratch every time. Treat a good prompt like a good template. The model changes; a well-crafted prompt holds up.

Ready to go deeper on one of the most powerful techniques in the toolkit? In the next chapter, we'll look at persona-based prompting — how assigning a role to the model changes not just the tone, but the quality and focus of everything it produces.
